

AUTONOMIS

• ENTERPRISE AI SOLUTIONS

AI Agents for Enterprise

Built for Production, Not for Demos

We deploy scalable, observable, and governed AI agents,
ready to operate within real business processes.

 Built-in Evaluation

 Full Traceability

 Control & Compliance

- THE CHALLENGE

Why Most AI Projects Never Reach Production

Enterprises invest heavily in AI experimentation, but few projects survive the journey from POC to production. The gap between "it works in a demo" and "it runs reliably at scale" is where most initiatives fail.

- ✖ **No Observability:** Black-box LLM calls make debugging impossible and costs unpredictable
- ✖ **Missing Guardrails:** No protection against hallucinations, prompt injection, or PII exposure
- ✖ **Unvalidated Outputs:** No systematic way to ensure response quality before customer delivery
- ✖ **Vendor Lock-in:** Tightly coupled architectures that can't adapt as better models emerge

- OUR APPROACH

The DOE Framework

A proven methodology for deploying AI agents that are governed from day one, fully observable, and built on modular architecture that grows with your needs.

D

Directives

Markdown-driven behavior

O

Orchestration

Multi-agent coordination

E

Execution

Production infrastructure

Enterprise-Grade Agent Architecture

Every component designed for reliability, observability, and governance at scale.



Autonomous Orchestration

Multi-agent teams (Triage → Response → Escalation) that resolve complex cases without human intervention.



Safety Guardrails

Block prompt injections, sanitize PII, and validate every response before it reaches customers.



LLM-as-Judge Evaluation

Automatic quality verification: relevance, tone, completeness, and actionability scoring on every output.



Complete Observability

Full pipeline traces: execution times, tokens consumed, costs, and performance metrics in real-time.



Modular Architecture

Swap LLMs (Gemini, OpenAI, Claude), CRMs, or evaluation metrics without changing business logic.



Directive-Driven Behavior

Non-engineers can edit agent behavior through simple markdown files—no code changes required.

BUILT WITH INDUSTRY-LEADING TOOLS

LangGraph

CrewAI

Langfuse

DeepEval

Guardrails AI

• IMMEDIATE AVAILABILITY

Let's Talk About Your AI Strategy

Free 30-Minute Strategy Call

-  **Needs Analysis** — Evaluate AI fit for your processes
-  **Process Audit** — Map and identify automation opportunities
-  **Recommendations** — Clear next steps tailored to your goals

 Book Your Call

cal.com/bernardo-garza-nlu8ko/30-min-meeting-english

Bernardo Garza

Founder, Autonomis

 autonomis.co

 bernardo@autonomis.co